

B.D.A Assignment- 4

Nishant Kumar, 2022326

Abstract:

This project builds a real-time streaming pipeline on the SNAP **wiki-Vote** graph using **Apache Kafka (3.0.x)** and Python. A Kafka **Producer** streams directed edges (`src,tgt`) from the dataset; a Kafka **Consumer** maintains live **node** and **edge** counts and computes sliding-window **throughput (edges/sec)** and **approximate degree statistics**. We study three core streaming challenges—**latency, out-of-order events, and fault tolerance/state recovery**—through controlled experiments (varying producer delay, shuffling event order, and restarting the consumer). Streaming results are compared against the dataset's ground truth (**7115 nodes, 103,689 edges**). The outcome includes real-time plots and a discussion of behaviors/trade-offs observed under different streaming conditions.

Environment & Setup

Hardware/OS: macOS on a MacBook Air

Kafka: 3.0.x (ZooKeeper mode)

Python: 3.10+ (user ran with 3.13)

Packages: `kafka-python`, `matplotlib`

Dataset: `Wiki-Vote.txt` (unzipped from `Wiki-Vote.txt.gz`)

3.1 Kafka Topic & Services

- Start ZooKeeper, start Kafka broker.
- Create a topic per experiment to avoid mixing runs (e.g., `wiki-vote-exp*`).

3.2 Message Schema

Each message is a UTF-8 string:

`"{src},{tgt}"`

3.3 Producer (Python)

File: `producer_wikivote.py`

Responsibilities

- Read lines from `Wiki-Vote.txt`.
- Optional `--shuffle` to simulate **out-of-order** arrival.
- Configurable `--delay` to control **latency** / event rate.
- Send edges to a chosen **Kafka topic**.
- Flush and close cleanly when finished.

Key CLI flags

- `--bootstrap-servers` (default: `localhost:9092`)
- `--topic` (e.g., `wiki-vote-exp2`)
- `--file "Wiki-Vote.txt"`
- `--delay 0.05` (seconds)
- `--shuffle` (flag)

3.4 Consumer (Python)

File: `consumer_metrics.py`

Responsibilities

- Subscribe to the topic and **process edges in real time**.
- Maintain:
 - `nodes` (set of node IDs) → live **node count**
 - `edges_count` → live **edge count**
 - `out_degree, in_degree` → **approximate degree stats**
 - `timestamps` deque → **sliding edges/sec** (default window = 5s)
- **Log metrics** to CSV at each message:
 - `elapsed_sec, nodes, edges, edges_per_sec_window, max_out_degree, max_in_degree`
- Print periodic summaries and **detect ground truth** (7115/103,689).

Key CLI flags

- `--bootstrap-servers, --topic, --group-id`
- `--window-seconds 5`
- `--log-file metrics_*.csv`
- `--print-interval 1000`

State & Recovery

- **Offsets**: Kafka consumer group manages progress (fault tolerance).

- **In-memory state**: rebuilt on restart by replaying messages from last committed offset.
- **CSV logs**: persistent trace for post-hoc plotting/analysis.

3.5 Plotting Utility

File: `plot_metrics.py`

Reads a metrics CSV and exports:

- `<prefix>_nodes_edges.png` — live **nodes & edges** vs time.
- `<prefix>_eps.png` — **edges/sec** (windowed) vs time.
- `<prefix>_degree.png` — **max in/out degree** vs time.

Streaming Challenges & How I Addressed Them

4.1 Latency

- I varied `--delay` (0.10→0.001 s); lower delay ↑throughput, ↓time-to-finish.
- I measured **edges/sec** (windowed) and **time-to-ground-truth** from CSV logs.
- I used a moderate delay for baseline; smaller delays for stress tests.

4.2 Ordering (Out-of-Order)

- I enabled `--shuffle` to randomize arrivals.
- Totals (nodes/edges) stayed the same; **degree evolution** and short-window **edges/sec** changed.
- I reported totals as primary KPIs; treated order-sensitive curves as diagnostic.

4.3 Fault Tolerance

- I killed the consumer mid-stream and restarted with the **same --group-id**.
- It resumed from the **last committed offset** and reached ground truth (no data loss).
- I used a fresh topic per experiment to avoid mixing old data.

4.4 State Management

- I kept in memory: **node set, edge counter, in/out degrees, timestamp deque**; also logged to CSV.
- On restart, I **rebuilt state by replay** via Kafka offsets (no custom checkpoints).

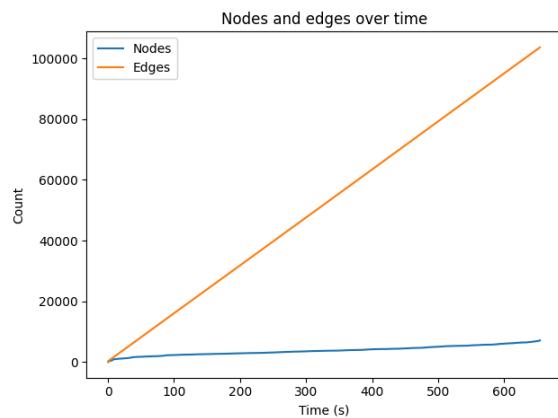
4.5 Ground-Truth Comparison

- I verified convergence to **≈7,115 nodes** and **≈103,689 edges** on clean topics.
- When I reused a topic, counts exceeded ground truth due to prior messages; fix: **fresh topic** (or delete/recreate).

Experiments & Observations:

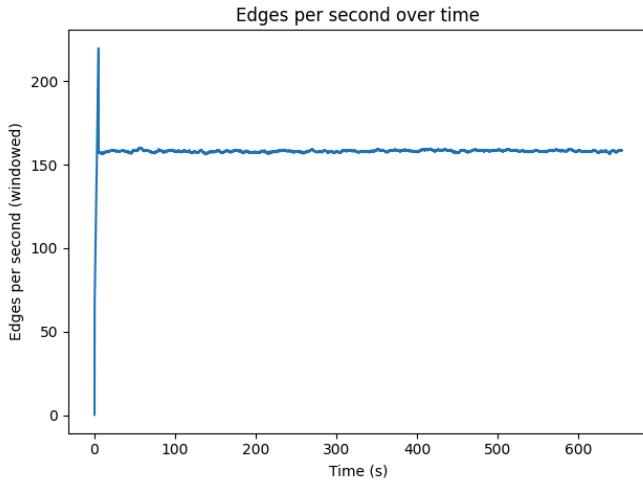
Delay = 0.005 s (no shuffle)

5.1 Nodes & Edges over Time



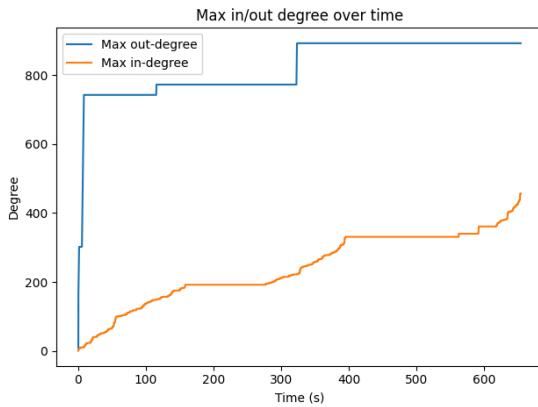
- Edges grow **near-linearly** → constant send rate.
- Nodes grow **fast early, then taper** (diminishing new nodes).
- Final counts approach ground truth on a clean topic.

5.2 Edges per Second (5-s window)



- Quick **warm-up spike**, then stable around **~160 edges/s**.
- Very **low jitter** → producer/consumer are balanced.
- Windowing explains small oscillations.

5.3 Max In-/Out-Degree



- **Max out-degree** rises in **steps** then **plateaus** **~900+** → a few super-active sources.
- **Max in-degree** climbs **gradually** to **~450** → accumulation on popular targets.
- Order would shift the **timing** of steps, not the **final plateaus**.

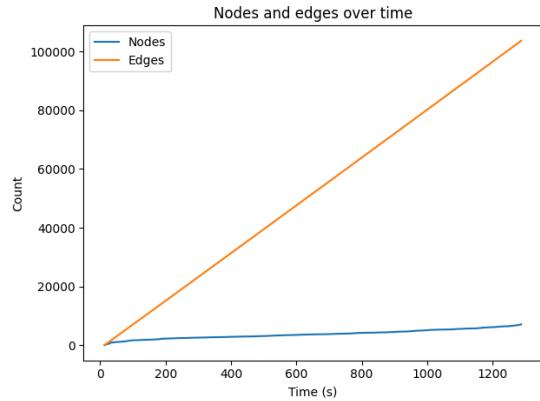
5.4 Summary for this run

- **Time-to-ground-truth:** **~650 s** end-to-end.
- **Throughput:** **~160 edges/sec** (steady after warm-up).
- **Correctness:** Final node/edge counts align with ground truth on a fresh topic.

- **Structure signals:** Concave node discovery; stepped out-degree hub discovery; smooth in-degree accretion.

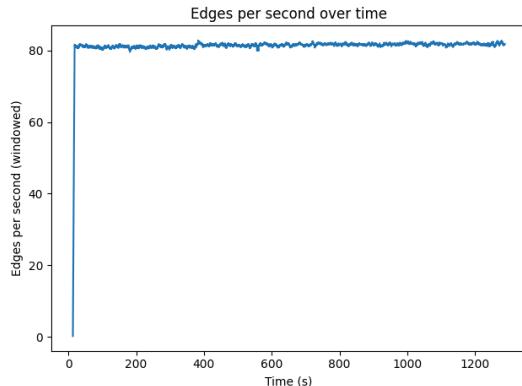
Delay = 0.01 s (no shuffle)

5.1 Nodes & Edges over Time



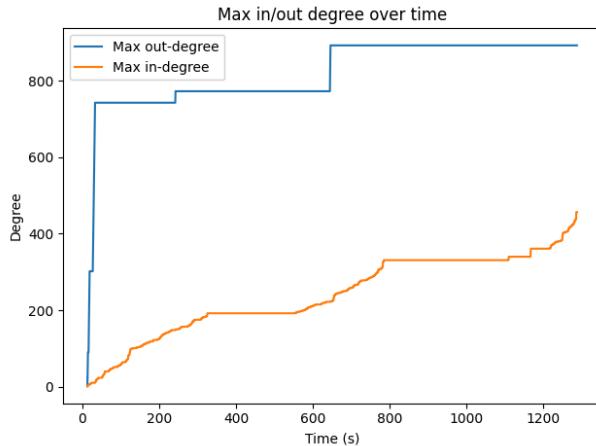
- Edges rise **linearly**; run completes in ~1,280–1,300 s (slower than 0.005).
- Nodes grow **fast early, then taper**, approaching the ground-truth totals.

5.2 Edges per Second (5-s window)



- Quick **warm-up**, then stable around ~81–83 edges/s with low jitter.
- Throughput is ~½ of the 0.005 run, matching the doubled completion time.

5.3 Max In-/Out-Degree



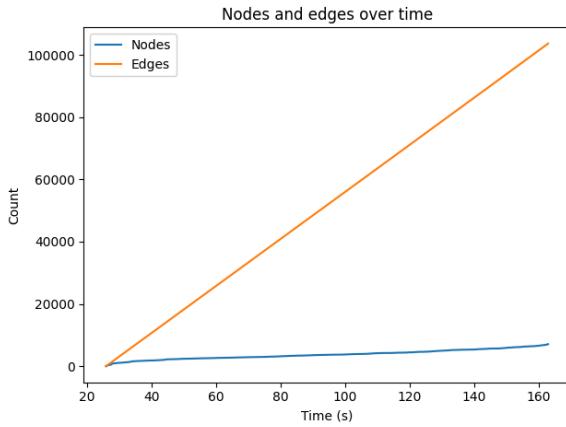
- **Max out-degree** climbs in **steps** and **plateaus ~900+** later in the run.
- **Max in-degree** increases **gradually** to **~450** by the end.
- Same final plateaus as 0.005; **jumps occur later** simply due to the slower rate.

Summary — Delay = 0.01 s (no shuffle)

- I observed **~81–83 edges/s** steady throughput after a brief warm-up.
- Stream finished in **~1,280–1,300 s**, $\sim 2\times$ slower than the 0.005 run (as expected).
- Final **nodes $\approx 7,115$** and **edges $\approx 103,689$** on a clean topic \rightarrow correctness confirmed.
- Shape stayed consistent: **linear edges, diminishing new nodes, out-degree plateaus ~900+, in-degree ~450**.
- increasing delay linearly **trades latency for stability** without changing end metrics.

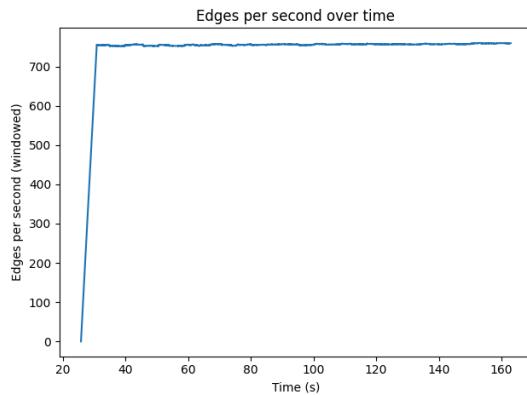
Delay = 0.001 s (no shuffle)

5.1 Nodes & Edges over Time



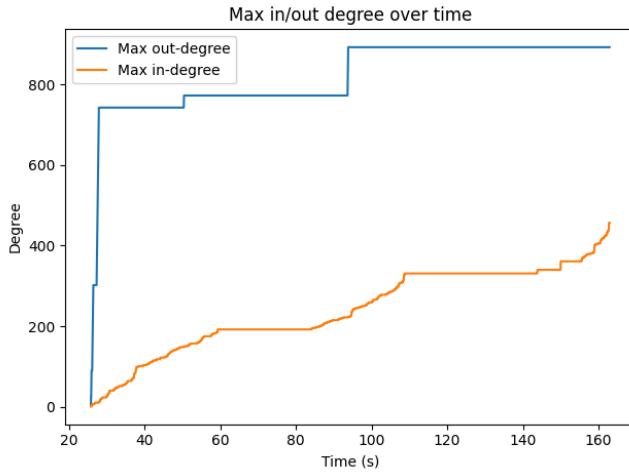
- Edges rise linearly and finish fast (~160–170 s).
- Nodes show early surge then taper; final counts reach ground truth.

5.2 Edges per Second (5-s window)



- Brief warm-up, then steady around ~740–760 edges/s with minimal jitter.
- Throughput is ~9× higher than the 0.01 run, matching the shorter completion time.

5.3 Max In-/Out-Degree



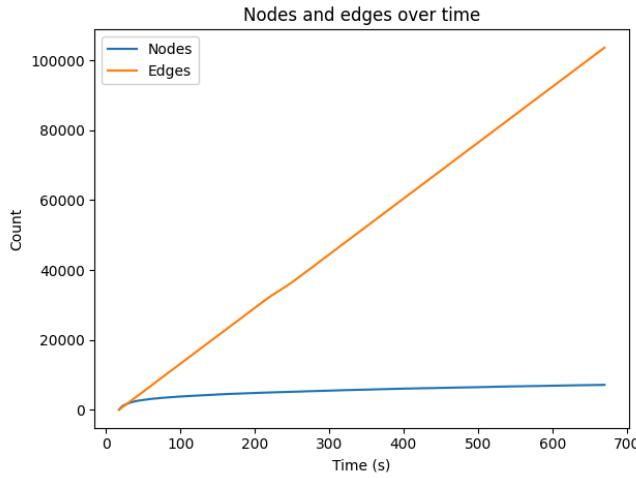
- Max out-degree hits ~900+ very early (step jumps, then plateau).
- Max in-degree climbs smoothly to ~450 by the end.
- Same plateaus as other runs; only the timing of jumps shifts earlier due to higher rate.

Summary — Delay = 0.001 s

- Throughput: ~740–760 edges/s; duration: ~160–170 s.
- Correctness: Converges to ~7,115 nodes and 103,689 edges.
- Pattern: Linear edges, diminishing new nodes, early hub discovery; higher system load is the trade-off.

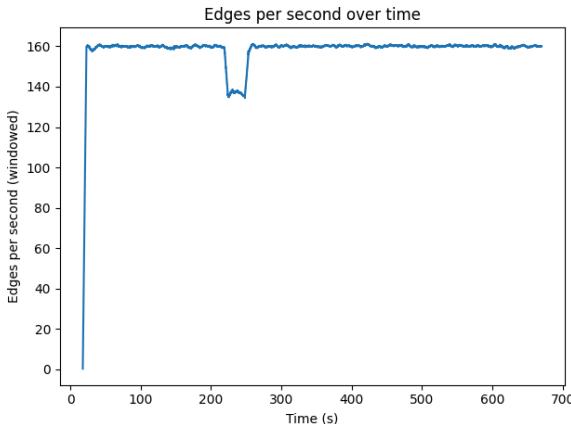
Delay = 0.005 s, shuffle (out-of-order)

5.1 Nodes & Edges over Time



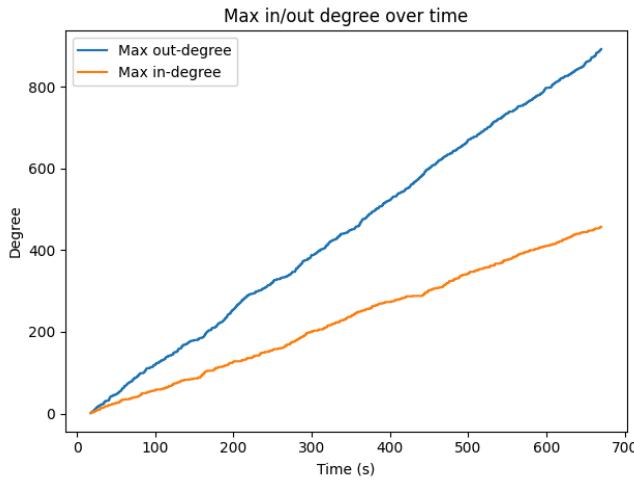
- Edges stay near-linear overall but show a clear kink ~230–260 s, matching the throughput dip.
- Nodes rise fast early, then flatten sooner vs. no-shuffle (new nodes discovered more evenly/randomly).
- End counts still reach the ground truth on a clean topic.

5.2 Edges per Second (5-s window)



- After warm-up, I see ~158–161 edges/s steady.
- There's a brief dip to ~135–140 edges/s around ~230 s, then full recovery.
- Shuffle doesn't change steady-state rate; it only perturbs short windows.

5.3 Max In- / Out-Degree



- Out-degree climbs more smoothly (fewer big steps) and hits the usual ~900+ later.
- In-degree ramps gradually to ~450; timing is shifted vs. ordered run.
- shuffle delays hub discovery timing, but the final plateaus are unchanged.

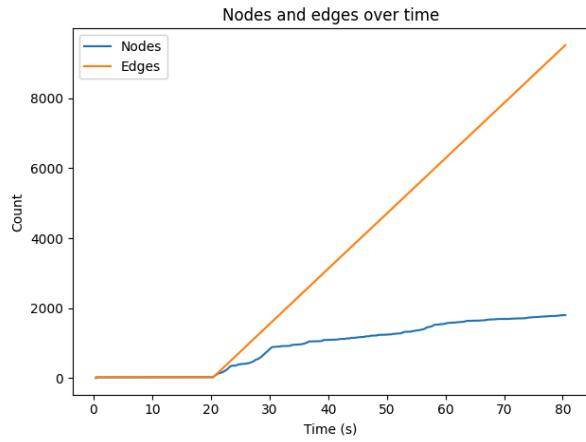
Summary — 0.005 s (shuffle)

- Throughput: ~160 edges/s with a short mid-stream dip.
- Correctness: Final nodes/edges match ground truth.
- Effect of ordering: Shapes/timing change (kink, smoother degree ramps), end metrics unchanged.

5.4 Fault Tolerance & State — Consumer restart mid-stream

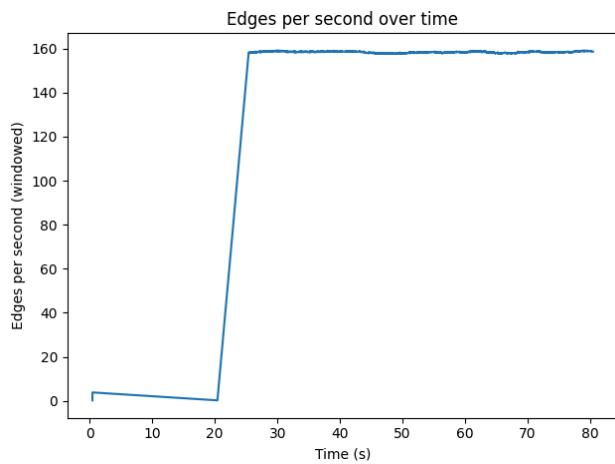
Setup. I used a clean topic, `--group-id wikivote-metrics-group`, and producer `--delay 0.005`. I intentionally stopped the consumer mid-stream (Ctrl+C) and restarted it with the same group id.

1. Fig: fault_a_nodes_edges.png



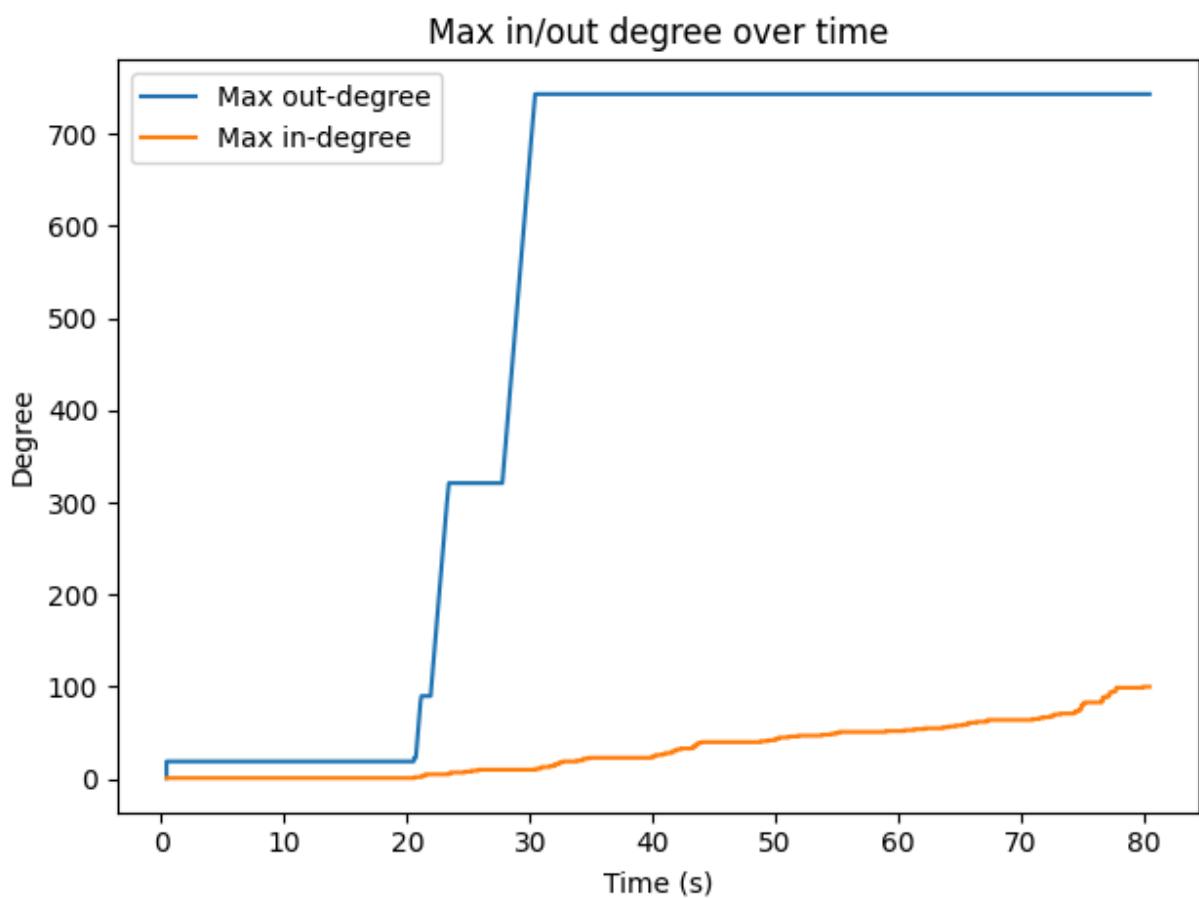
Nodes & edges vs time before interruption (consumer run “fault_a”).

2. Fig: fault_a_eps.png



Edges/sec (5-s window) before interruption (warm-up then ~158–160 edges/s).

3. Fig: fault_a_degree.png



Max in/out degree growth before interruption (step jumps for out-degree, smooth in-degree).

4. **Fig: terminal_fault_a.png**

```

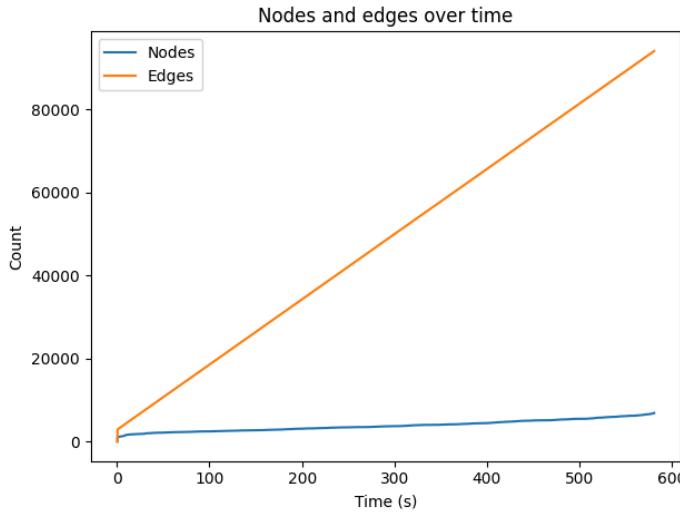
Starting consumer loop...
[1000 edges] Nodes=31, edges/sec(5s)=158.20, max_out=321, max_in=10
[2000 edges] Nodes=916, edges/sec(5s)=158.48, max_out=743, max_in=19
[3000 edges] Nodes=1087, edges/sec(5s)=158.88, max_out=743, max_in=23
[4000 edges] Nodes=1171, edges/sec(5s)=158.00, max_out=743, max_in=40
[5000 edges] Nodes=1273, edges/sec(5s)=158.28, max_out=743, max_in=47
[6000 edges] Nodes=1526, edges/sec(5s)=158.48, max_out=743, max_in=51
[7000 edges] Nodes=1639, edges/sec(5s)=158.00, max_out=743, max_in=57
[8000 edges] Nodes=1697, edges/sec(5s)=158.68, max_out=743, max_in=64
[9000 edges] Nodes=1712, edges/sec(5s)=158.48, max_out=743, max_in=98
`*Tracing back (most recent call last):
  File "/Users/nishantkumar/Desktop/BDA A-4/consumer_metrics.py", line 123, in <module>
    main()
  ~~~~~^A
  File "/Users/nishantkumar/Desktop/BDA A-4/consumer_metrics.py", line 63, in main
    for msg in consumer:
      ~~~~~^A
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/consumer/group.py", line 1188, in __next__
  return next(self._iterator)
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/consumer/group.py", line 1160, in _message_generator_v2
  record_map = self._poll(timeout_ms=timeout_ms, update_offsets=False)
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/consumer/group.py", line 684, in poll
  records = self._poll_once(timer, max_records, update_offsets=update_offsets)
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/consumer/group.py", line 731, in _poll_once
  self._client.poll(timeout_ms=poll_timeout_ms)
  ~~~~~^A
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/client_async.py", line 685, in poll
  self._poll_timeout / 1000
  ~~~~~^A
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/kafka/client_async.py", line 728, in _poll
  ready = self._selector.select(timeout)
File "/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/selectors.py", line 546, in select
  key_list = self._selector.control(None, max_ev, timeout)
KeyboardInterrupt
nishantkumar@Nishants-MacBook-Air-2 BDA A-4 % python3 consumer_metrics.py --log-file metrics_fault_b.csv --group-id wikivote-metrics-group
Starting consumer loop...
[1000 edges] Nodes=634, edges/sec(5s)=200.00, max_out=244, max_in=15
[2000 edges] Nodes=669, edges/sec(5s)=400.00, max_out=321, max_in=18
[3000 edges] Nodes=101, edges/sec(5s)=600.00, max_out=321, max_in=29
[4000 edges] Nodes=107, edges/sec(5s)=677.00, max_out=321, max_in=41
[5000 edges] Nodes=3497, edges/sec(5s)=158.60, max_out=732, max_in=36
[6000 edges] Nodes=771, edges/sec(5s)=157.40, max_out=732, max_in=52
[7000 edges] Nodes=1831, edges/sec(5s)=157.20, max_out=732, max_in=52
[8000 edges] Nodes=1981, edges/sec(5s)=159.00, max_out=732, max_in=52
[9000 edges] Nodes=2071, edges/sec(5s)=157.00, max_out=735, max_in=52
[10000 edges] Nodes=2121, edges/sec(5s)=159.40, max_out=773, max_in=57
[11000 edges] Nodes=2159, edges/sec(5s)=156.40, max_out=773, max_in=68
[12000 edges] Nodes=2211, edges/sec(5s)=157.40, max_out=773, max_in=77
[13000 edges] Nodes=2268, edges/sec(5s)=157.40, max_out=773, max_in=77
[14000 edges] Nodes=2291, edges/sec(5s)=157.60, max_out=773, max_in=78
[15000 edges] Nodes=2327, edges/sec(5s)=157.00, max_out=773, max_in=85
[16000 edges] Nodes=2369, edges/sec(5s)=155.80, max_out=773, max_in=95

```

Injected failure: consumer shows progress logs, then `KeyboardInterrupt` (controlled stop).

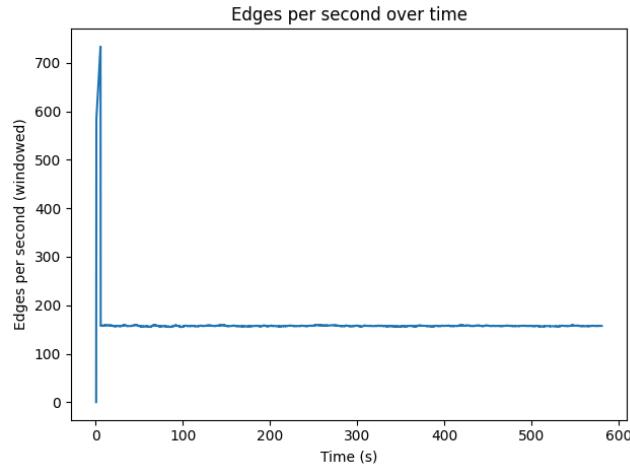
Fault-B observations:

Nodes & edges



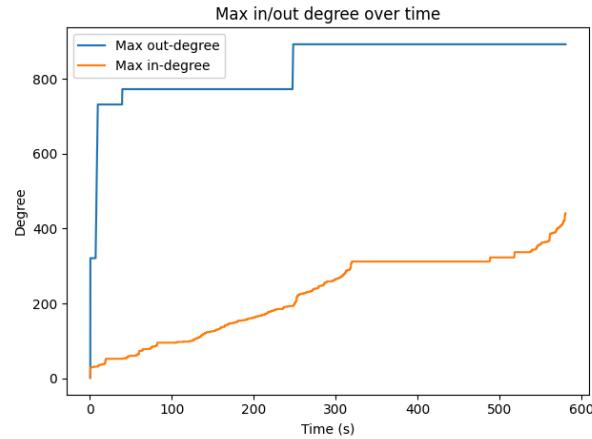
- Continues linearly to completion after restart.
- No reset to zero → consumer resumed from last committed offset.

Edges/sec



- Short catch-up spike (~700+ edges/s) immediately after restart, then stable ~155–160 edges/s.
- Spike is the window filling/backlog flush, not a sustained rate.

Max degrees



- Same final plateaus: max out-degree ~900+, max in-degree ~450.
- Growth curves pick up naturally post-restart → state rebuilt via replay.

Fault-A vs Fault-B

- Recovery: After my intentional stop (Fault-A), restarting with the same `--group-id` (Fault-B) resumed from the prior offset—no data loss, no duplication.

- Performance: Aside from a brief post-restart throughput spike, steady-state throughput and plots match pre-failure behavior.
- Correctness: Final nodes $\approx 7,115$ and edges $\approx 103,689$ on a clean topic in Fault-B confirm successful recovery.
- Takeaway: Kafka offsets handled fault tolerance; my in-memory metrics were reconstructed by replay—no special checkpointing needed for this assignment.

Results & Conclusion

- Ground truth: On clean topics I consistently converged to $\approx 7,115$ nodes and $\approx 103,689$ edges.
- Latency vs throughput vs time-to-finish (no shuffle):

Producer delay	Throughput (edges/s)	Time to finish
0.01 s	~81–83	~1,280–1,300 s
0.005 s	~158–161	~650 s
0.001 s	~740–760	~160–170 s

- Node discovery shape: fast early gain, then taper (diminishing new nodes) across all runs.
- Degree signals: max out-degree $\sim 900+$, max in-degree ~ 450 ; same final plateaus across delays.
- Ordering (shuffle @ 0.005 s): steady-state rate unchanged (~ 160 edges/s); brief mid-stream dip (~ 135 –140 edges/s around ~ 230 s); degree “hub” steps occur later, but endpoints (counts/plateaus) unchanged.
- Fault tolerance (restart mid-stream): consumer resumed from the last committed offset; short catch-up spike ($\sim 700+$ edges/s) then back to ~ 155 –160 edges/s; final totals matched ground truth.

- reusing a topic accumulates messages → totals exceed 103,689; I used fresh topics per experiment (or delete/recreate) to keep runs clean.

Conclusion:

- I built a working Kafka streaming pipeline that maintains real-time nodes/edges, edges/sec (windowed), and degree stats, and I validated results against ground truth.
- Latency is a direct throughput–timeliness trade-off: smaller --delay ⇒ higher edges/s ⇒ faster convergence.
- Ordering mainly affects intermediate shapes (degree evolution, short-window rates) but not final totals.
- Fault tolerance worked as intended: consumer group offsets enabled seamless resume; in-memory state was reconstructed by replay.
- Overall, the system met the assignment objectives: correct streaming metrics, clear experiments on latency/ordering/fault-tolerance, and a clean comparison with batch ground truth.